



Big Data in NYC for Taxi rides management with Apache Hive

Gabriele Favia matr. 579166

December 2019

Contents

The NYC Taxi and Limousine Commission.....	3
Requests	4
Data dictionary.....	5
DDL and DML commands	6
DQL commands	10
Dealing with different datasets	11
Dealing with different operational modes.....	13
Results	14
Performance comparison.....	22
Limitations	24
References	24
Appendix	25

The NYC Taxi and Limousine Commission

The New York City Taxi and Limousine Commission (NYC TLC) is an agency of the New York City government that licenses and regulates the medallion taxis and for-hire vehicle industries, including app-based companies.

The TLC's regulatory landscape includes medallion (yellow) taxicabs, green or Boro taxicabs, black cars (including both traditional and app-based services), community-based livery cars, commuter vans, paratransit vehicles (ambulettes), and some luxury limousines [1].


But what is the difference between yellow and green taxi (cab in American English)?

The famous NYC yellow taxis provide transportation exclusively through street-hails.

The number of taxicabs is limited by a finite number of medallions issued by the TLC.

It's possible to access this mode of transportation by standing in the street and hailing an available taxi with your hand. The pickups are not pre-arranged.

The green taxis know as well as “Boro” cab or Street Hail Livery (SHL) are permitted to accept street-hails above 110th Street in Manhattan and in the outer-boroughs of New York City. The SHL program allows livery vehicle owners to license and outfit their vehicles with green borough taxi branding, meters, credit card machines, and ultimately the right to accept street hails in addition to pre-arranged rides [2].



Requests

Designing and implement in Hadoop MapReduce:

4. A HiveQL query returning the total number of passengers in yellow and green taxi in 2018.
5. A HiveQL query returning, for each rate code, the total amount charged to passengers of yellow and green taxi in 2018.
6. A HiveQL query returning, for each *PULocationID*, the list of related *DOLocationID* for yellow taxi in 2018, ordered by increasing average trip distance.

Data dictionary

The two kind of file of interest, yellow and green have a different header, which means that the software has the need to distinguish and operate the two cases differently.

Moreover, the green and yellow taxis dataset produced from January 2016 to June 2016, follow a different data dictionary that requires some variation to the parsing procedure for the points 4 and 5, while for the point 6, the dataset cannot be use because of the absence of these fields.

The fields of interest to be extracted from both the yellow and green cabs files are:

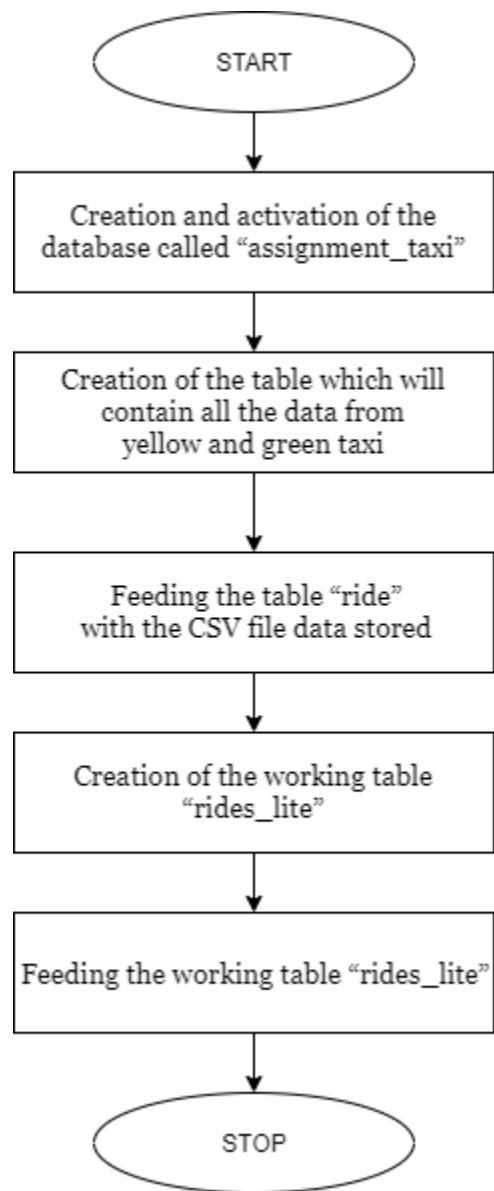
Field	Description
lpep_pickup_datetime or rpep_pickup_datetime	The pickup time, useful to get the year.
Passenger_count	The number of passengers in the vehicle.
RatecodeID	The final rate code in effect at the end of the trip. 1 = Standard rate 2 = JFK 3 = Newark 4 = Nassau or Westchester 5 = Negotiated fare 6 = Group ride
total_amount	The total amount charged to passengers. Does not include cash tips.

While the fields of interest exclusively extracted from yellow files are:

Field	Description
PULocationID	TLC Taxi Zone in which the taximeter was engaged.
DOLocationID	TLC Taxi Zone in which the taximeter was disengaged.
trip_distance	The elapsed trip distance in miles reported by the taximeter.

A comprehensive description of all the fields is available at nyc.gov website both for yellow [3] and the green [4] cabs.

DDL and DML commands



Creation and activation of the database called “assignment_taxi”

```
CREATE DATABASE IF NOT EXISTS assignment_taxi;  
USE assignment_taxi;
```

Creation of the table which will contain all the data from yellow and green taxi

```
CREATE TABLE rides (  
VendorID INT,  
lpep_pickup_datetime STRING,  
lpep_dropoff_datetime STRING,  
store_and_fwd_flag STRING,  
RatecodeID INT,  
PULocationID INT,  
DOLocationID INT,  
passenger_count INT,  
trip_distance FLOAT,  
fare_amount INT,  
extra INT,  
mta_tax INT,  
tip_amount INT,  
tolls_amount INT,  
ehail_fee STRING,  
improvement_surcharge INT,  
total_amount INT,  
payment_type INT,  
trip_type INT  
) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' tblproperties  
("skip.header.line.count"="2");
```

- ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' sets the delimiter of the columns to be a comma.
- tblproperties ("skip.header.line.count"="2") makes the process of data importing skipping the first two rows of the CSV document, which are made up of a header and a blank row.

Feeding the table “ride” with the CSV file data stored

```
LOAD DATA LOCAL INPATH '/media/sf_condiv/Hive_7Dec/input' OVERWRITE INTO  
TABLE rides;
```

Creation of the working table “rides_lite”

To make the queries faster and reducing the memory usage at the same time, a smaller version of the loaded table is created, containing only the fields useful for the queries (the ones described in [Data dictionary](#) section).

```
CREATE TABLE rides_lite (  
year INT,  
RatecodeID INT,  
PULocationID INT,  
DOLocationID INT,  
passenger_count INT,  
trip_distance FLOAT,  
total_amount INT,  
taxi_type INT  
);
```

- The `taxi_type` field is added with the aim of distinguishing with ease (just a INT flag) whether the row of interest is relative to a green (value 1) or a yellow cab (value 2).

Feeding the working table “rides_lite”

```
INSERT OVERWRITE TABLE rides_lite  
SELECT  
YEAR(lpep_pickup_datetime) as year,  
RatecodeID,  
PULocationID,  
DOLocationID,  
passenger_count,  
trip_distance,  
total_amount,  
CASE WHEN (REVERSE(SPLIT(REVERSE(INPUT__FILE__NAME), '[/]')[0])) LIKE  
'green%' THEN 1 WHEN (REVERSE(SPLIT(REVERSE(INPUT__FILE__NAME),  
'[/]')[0])) LIKE 'yellow%' THEN 2 END  
FROM rides;
```

- Noticing that the objective queries require only the year of a ride, it's been chosen to insert only this inside the `rides_lite` table using the function `YEAR` which extracts the year from a string date. This expedient lightens the computational and memory requirements, also because the year is here treated as an integer instead of a string.
- To value the field `taxi_type` of the table `rides_lite`, Hive checks if the file name starts with the keyword green or yellow:
The keyword `INPUT__FILE__NAME` is a virtual column [5] which contains the full path of the source file related to that row.
The algorithm consists of reversing the full path and getting the first word delimited by a slash (resulting to be the file name with reversed letters disposition), extracting the first element and reversing it again to get the file name of the CSV file; then, depending by the type of the taxi, Hive assigns 1 or 2 to the `taxi_type` field using the `CASE` statement and the `LIKE` operator.

INPUT__FILE__NAME	hdfs://localhost:9000/user/hive/warehouse/assignment_taxi.db/rides/green_tripdata_2016-07.csv
REVERSE	vsc.70-6102_atadpirt_neerg/sedir/bd.ixat_tnemngissa/esuoheraw/evih/resu/0009:tsohlacol//:sfdh
SPLIT on /	[0 => vsc.70-6102_atadpirt_neerg] [1 => sedir] [2 => bd.ixat_tnemngissa] [3 => esuoheraw] [4 => evih] [5 => resu] [6 => 0009:tsohlacol] [7 =>] [8 => :sfdh]
SPLIT_ELEMENT[0]	vsc.70-6102_atadpirt_neerg
REVERSE	green_tripdata_2016-07.csv

DQL commands

To store the result in a file it is possible to put before the query, the following command:

```
INSERT OVERWRITE LOCAL DIRECTORY 'home/result_hive'
```

Query 4

```
SELECT year, SUM(passenger_count) FROM rides_lite GROUP BY year HAVING  
year = '2018';
```

- This query uses the SUM function to sum the *passenger_count* and group the results by year.

Query 5

```
SELECT RatecodeID, SUM(total_amount) FROM rides_lite WHERE year = '2018'  
GROUP BY RatecodeID;
```

This query uses the SUM function to sum the *total_amount*, for rides in 2018 and group the results by RatecodeID.

Query 6

```
SELECT PULocationID, COLLECT_SET(DOLocationID) as group_doloc,  
COLLECT_LIST(avg_trip) as group_trip FROM (SELECT PULocationID,  
DOLocationID, CAST(FORMAT_NUMBER(AVG(trip_distance), 2) AS FLOAT) AS  
avg_trip FROM rides_lite WHERE taxi_type = 1 AND year = '2018' GROUP BY  
PULocationID, DOLocationID ORDER BY avg_trip) table_temp GROUP BY  
PULocationID;
```

This query is made up of two queries:

- Internal query: selecting all the yellow taxi (*taxi_type* = 1), it gets the values of the columns *PULocationID* and *DOLocationID* and groups the average of *trip_distance* by *PULocationID*. The average results are expressed in miles, approximated to the second decimal digit, for ease of reading purposes. The results are sort by average trip.
- External query: groups by *PULocationID* and display the results of *DOLocationID* and *avg_trip* in a “list-like” mode.

Dealing with different datasets

To make the measurement comparison three datasets have been created:

- A table containing data for 2018, 2017 e the second half of 2016, inside the table *rides_lite*.
- A table containing data for 2018 and 2017 only, called *rides_lite_2018_2017*.
- A table containing data from 2018 only, called *rides_lite_2018*.

Construction and population of the table for the years 2018 and 2017

```
CREATE TABLE rides_lite_2018_2017 (  
  year INT,  
  RatecodeID INT,  
  PULocationID INT,  
  DOLocationID INT,  
  passenger_count INT,  
  trip_distance FLOAT,  
  total_amount INT,  
  taxi_type INT  
);  
  
INSERT OVERWRITE TABLE rides_lite_2018_2017  
SELECT  
  YEAR(lpep_pickup_datetime) as year,  
  RatecodeID,  
  PULocationID,  
  DOLocationID,  
  passenger_count,  
  trip_distance,  
  total_amount,  
  CASE WHEN (REVERSE(SPLIT(REVERSE(INPUT__FILE__NAME), '[/]')[0])) LIKE  
    'green%' THEN 1 WHEN (REVERSE(SPLIT(REVERSE(INPUT__FILE__NAME),  
    '[/]')[0])) LIKE 'yellow%' THEN 2 END  
FROM rides WHERE lpep_pickup_datetime LIKE '2018%' OR lpep_pickup_datetime  
LIKE '2017%';
```

Construction and population of the table for the year 2018

```
CREATE TABLE rides_lite_2018 (  
  year INT,  
  RatecodeID INT,  
  PULocationID INT,  
  DOLocationID INT,  
  passenger_count INT,  
  trip_distance FLOAT,  
  total_amount INT,  
  taxi_type INT  
);  
  
INSERT OVERWRITE TABLE rides_lite_2018
```

```

SELECT
YEAR(lpep_pickup_datetime) as year,
RatecodeID,
PULocationID,
DOLocationID,
passenger_count,
trip_distance,
total_amount,
CASE WHEN (REVERSE(SPLIT(REVERSE(INPUT__FILE__NAME), '['/']')[0])) LIKE
'green%' THEN 1 WHEN (REVERSE(SPLIT(REVERSE(INPUT__FILE__NAME),
 '['/']')[0])) LIKE 'yellow%' THEN 2 END
FROM rides WHERE lpep_pickup_datetime LIKE '2018%';

```

Query 4 for the table rides_lite_2018_2017

```

SELECT year, SUM(passenger_count) FROM rides_lite_2018_2017 GROUP BY year
HAVING year = '2018';

```

Query 5 for the table rides_lite_2018_2017

```

SELECT RatecodeID, SUM(total_amount)FROM rides_lite_2018_2017 WHERE year =
'2018' GROUP BY RatecodeID;

```

Query 6 for the table rides_lite_2018_2017

```

SELECT PULocationID, COLLECT_SET(DOLocationID) as group_doloc,
COLLECT_LIST(avg_trip) as group_trip FROM (SELECT PULocationID,
DOLocationID, CAST(FORMAT_NUMBER(AVG(trip_distance), 2) AS FLOAT) AS
avg_trip FROM rides_lite_2018_2017 WHERE taxi_type = 1 AND year = '2018'
GROUP BY PULocationID, DOLocationID ORDER BY avg_trip) table_temp GROUP BY
PULocationID;

```

Query 4 for the table rides_lite_2018

```

SELECT year, SUM(passenger_count) FROM rides_lite_2018 GROUP BY year;

```

Query 5 for the table rides_lite_2018

```

SELECT RatecodeID, SUM(total_amount)FROM rides_lite_2018 GROUP BY
RatecodeID;

```

Query 6 for the table rides_lite_2018

```

SELECT PULocationID, COLLECT_SET(DOLocationID) as group_doloc,
COLLECT_LIST(avg_trip) as group_trip FROM (SELECT PULocationID,
DOLocationID, CAST(FORMAT_NUMBER(AVG(trip_distance), 2) AS FLOAT) AS

```

```
avg_trip FROM rides_lite_2018 WHERE taxi_type = 1 GROUP BY PULocationID,  
DOLocationID ORDER BY avg_trip) table_temp GROUP BY PULocationID;
```

Dealing with different operational modes

Another parameter to take in account when it comes to performance analysis is the use of the hive operational modes: local (exectution on the local machine as a simple node) and pseudo-distributed (simulation of a node in a network)

Activation of the local mode

```
SET mapred.job.tracker = local;  
SET hive.exec.mode.local.auto = true;
```

Activation of the pseudo-distributed node

```
SET mapred.job.tracker = pseudo;  
SET hive.exec.mode.local.auto = false;
```

Results

Query 4

Year	SUM(passenger_count)
2018	16789109123

Query 5

RatecodeID	SUM(total_amount)
3	135431705
35	393328
67	16894
99	4002
131	601
163	498
195	1433
547	0
611	44
18	38855698
50	87956
82	5783
114	2455
146	350
178	493
210	486
274	811
6098	10000
10	48361500
42	151954
74	10863
106	1746
138	466
20	16510505
52	78108

84	5190
116	1195
148	593
212	240
2	228543473
34	373812
66	28521
98	3858
130	1291
162	300
226	700
258	10005
290	729
6	49935919
38	257534
70	18206
102	1438
134	456
166	190
198	400
230	52
422	1059
25	1728785
57	42296
89	5796
121	931
153	1047
4	87835275
36	442198
68	18582
100	3300
132	1498
164	1
196	983

484	2417
932	14
1764	4414
12	24786499
44	135621
76	9504
108	685
140	763
204	1598
300	7
24	1983564
56	42497
88	3436
120	3028
152	277
216	302
248	412
600	5
5	70676335
37	344676
69	20014
101	10318
133	311
229	450
325	50
1061	8
2053	5134
5381	10001
22	5743867
54	59485
86	4615
118	3398
150	275
182	25

214	471
246	1322
31	584971
63	27367
95	3141
127	1363
159	782
831	11
7	41743771
39	225538
71	16419
103	3766
135	635
7655	10001
1	490008478
33	418931
65	27326
97	3820
129	2854
801	12
833	23
28	1648840
60	35305
92	4264
124	4166
156	66
188	10
252	1340
380	0
6204	10000
16	30579674
48	90516
80	4983
112	1801

144	510
176	2658
528	1325
4016	10000
19	23170653
51	75338
83	3399
115	1923
147	1051
211	1227
1459	3652
1811	4530
26	1874646
58	39006
90	4830
122	1150
154	462
218	338
250	656
602	18
1402	3509
1818	4551
17	49407742
49	90402
81	6218
113	2657
145	700
753	1887
1297	3247
1841	4606
0	206982960
32	539941
64	27540
96	5126

128	1353
27	1720835
59	36994
91	4318
123	1727
155	300
251	1177
603	26
1723	4311
14	15348514
46	101087
78	7740
110	2605
142	2405
206	581
270	25
302	234
910	300
23	2897899
55	56069
87	3805
119	2192
151	1822
183	0
2231	5581
9	52519319
41	172225
73	10075
105	1912
137	1021
169	372
329	0
13	16976114
45	119482

77	8688
109	3146
141	374
301	25
11	37619022
43	137584
75	6425
107	2154
139	1077
267	950
189483	4
29	1044709
61	33477
93	3641
125	1484
157	225
189	170
15	18439856
47	98777
79	8051
111	2853
207	502
239	639
943	14
21	11125701
53	61243
85	5157
117	3677
149	394
181	875
213	858
245	466
341	1002
8	47251101

40	188208
72	12932
104	2613
136	298
168	211
1928	4824
2408	6023
30	710537
62	30268
94	3368
126	1366
222	0
606	31
830	9

Query 6 (first row only)

PULocationID	COLLECT_SET(DOLocationID) AS group_doloc	COLLECT_LIST(avg_trip) AS group_trip
1	[1,265,23,249,66,68,14,164,22,243,145,166,257,64,132]	[1.06,8.93,12.7,12.86,15.95,17.1,18.71,20.1,22.0,23.9,25.21,26.64,32.9,34.9,37.1,161.19]

Performance comparison

The program instances are run on the following virtual machine:

The operative system is Ubuntu 18.4.4 hosted by VirtualBox vers. 6.0.14 which are given:

2 cores;
4096 MB RAM;
60GB disk space;

Equipped with Hadoop vers. 2.8.5, Java Open JDK vers. 1.8.0

The host machine with:

Intel® Core™ i7 6500U and SSD.

Time for completion

years: 2018
number of entries in table *rides_lite_2018*: 111609853




Task	Execution time in local mode (s)	Execution time in psuedo-distributed (s)
Query 4	118.137	131.078
Query 5	129.997	127.181
Query 6	134.667	147.617

years: 2017, 2018
number of entries in table *rides_lite_2018_2017*: 236847300

Task	Execution time in local mode (s)	Execution time in psuedo-distributed (s)
Query 4	199.514	178.525
Query 5	229.105	201.061
Query 6	157.725	161.184

years: 2016, 2017, 2018
number of entries in table *rides_lite*: 305975113

Task	Execution time in local mode (s)	Execution time in psuedo-distributed (s)
Query 4	354.845	211.382
Query 5	297.944	243.932
Query 6	409.915	525.028



The pseudo-distributed mode performs better or worse than the local mode, depending on the type of query and the amount of data: the biggest differences are noticeable in case of the whole dataset table (made up of over 300M rows) where the pseudo-distributed mode takes -40.42 % time in comparison to local mode in terms of completion time for the query 4, while it uses +28.08 % time more when the query 6 is run.

Limitations

The results concerning the execution times are biased by the state of the host machine which was loaded with different tasks in background, causing a not constant access time and bandwidth availability on the SSD on the host operating system, so these must be taken as an overall idea of the average execution time for the machine used.

References

- [1] https://en.wikipedia.org/wiki/New_York_City_Taxi_and_Limousine_Commission
- [2] <https://www.quora.com/What-is-the-difference-between-Green-Cabs-and-Yellow-Cabs>
- [3] https://www1.nyc.gov/assets/tlc/downloads/pdf/data_dictionary_trip_records_yellow.pdf
- [4] https://www1.nyc.gov/assets/tlc/downloads/pdf/data_dictionary_trip_records_green.pdf
- [5] <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+VirtualColumns>

Appendix

Screenshot of resulting file for Query 4

1	2018	16789109123
2	Time taken: 131.078 seconds, Fetched: 1 row(s)	

Screenshot of resulting file for Query 5

1	3	135431705	64	44	135621	126	89	4983	189	119	2192
2	35	3932328	65	76	9504	127	112	1801	190	151	1822
3	67	16894	66	108	685	128	144	510	191	183	0
4	99	4002	67	140	763	129	176	2658	192	2231	5581
5	131	601	68	204	1598	130	228	1325	193	0	52519319
6	163	498	69	300	7	131	4016	10000	194	41	172225
7	195	1433	70	24	1983564	132	19	23170653	195	73	10075
8	547	0	71	56	42497	133	51	75338	196	105	1912
9	611	44	72	88	3436	134	83	3399	197	137	1021
10	18	38855698	73	120	3028	135	115	1923	198	169	372
11	50	87956	74	152	277	136	147	1051	199	329	0
12	80	5783	75	216	302	137	211	1227	200	15	16976114
13	114	2455	76	248	412	138	1459	3652	201	45	119482
14	146	350	77	500	5	139	1811	4530	202	75	9688
15	178	493	78	5	70676335	140	26	1874646	203	109	3146
16	210	486	79	37	344676	141	58	39006	204	141	374
17	274	811	80	69	20014	142	90	4830	205	301	25
18	10	1698	81	101	16318	143	122	1150	206	11	37619822
19	10	48361500	82	133	311	144	154	462	207	43	137584
20	42	151954	83	229	450	145	218	338	208	75	6425
21	74	10863	84	235	59	146	250	656	209	107	2154
22	106	1746	85	1061	8	147	602	18	210	139	1077
23	138	466	86	2053	5134	148	1402	3509	211	267	950
24	20	16310505	87	5381	10001	149	1818	4561	212	189483	4
25	52	78108	88	22	5743867	150	17	49407742	213	29	1044709
26	84	5190	89	54	59485	151	49	90402	214	61	33477
27	116	1195	90	96	4615	152	81	6218	215	93	3641
28	148	593	91	118	3398	153	113	2637	216	125	1484
29	212	240	92	150	275	154	145	700	217	157	225
30	2	236543473	93	182	25	155	753	1887	218	189	170
31	34	373812	94	214	471	156	1297	3247	219	15	18439856
32	66	28521	95	246	1322	157	1841	4606	220	47	98777
33	98	3858	96	31	584971	158	0	206261960	221	79	8651
34	130	1291	97	63	27307	159	32	539941	222	111	2853
35	162	308	98	95	3141	160	64	27540	223	207	502
36	226	700	99	127	1363	161	96	5176	224	239	639
37	258	10095	100	159	782	162	128	1353	225	943	14
38	290	729	101	831	11	163	27	1728835	226	21	11125701
39	6	49935919	102	7	41743771	164	59	36994	227	53	61243
40	28	257534	103	29	225338	165	91	4310	228	85	5157
41	70	18206	104	71	16419	166	123	1727	229	117	3677
42	102	1438	105	103	3766	167	155	308	230	149	394
43	134	456	106	133	635	168	251	1177	231	181	875
44	166	190	107	7655	10001	169	603	26	232	213	858
45	198	409	108	1	490008478	170	1723	4311	233	245	466
46	230	52	109	33	418931	171	14	15348314	234	341	1002
47	422	1059	110	65	27326	172	46	101807	235	8	47251101
48	25	1728785	111	97	3820	173	78	7740	236	40	188208
49	57	42296	112	129	2854	174	110	2605	237	72	12932
50	89	5796	113	801	12	175	142	2405	238	104	2613
51	121	931	114	833	23	176	206	581	239	136	298
52	153	1047	115	28	1648840	177	270	25	240	168	211
53	4	67835275	116	60	35305	178	302	234	241	1926	4824
54	36	442198	117	92	4264	179	910	300	242	2408	6823
55	68	18582	118	124	4166	180	23	2897899	243	30	710537
56	100	3300	119	154	66	181	55	56669	244	62	38268
57	132	1498	120	188	10	182	87	3805	245	94	3368
58	164	1	121	252	1340	183	119	2192	246	126	1366
59	196	883	122	380	0	184	151	1822	247	222	0
60	484	2417	123	6284	10000	185	183	0	248	606	31
61	932	14	124	16	30579674	186	2231	5581	249	830	9
62	1784	4414	125	48	90516	187	9	52519319	250	Time taken: 243.932 seconds, Fetched: 243 row(s)	
63	12	24786499	126	80	4983	188	41	172225			

Screenshot of resulting file for the Query 6 (first row only)

1	[1,265,23,249,66,68,14,164,22,243,145,166,257,64,132]	[1.06,8.93,12.7,12.86,15.95,17.1,18.71,20.1,22.0,23.9,25.21,26.64,32.9,34.9,37.1,161.19]
---	---	--

Row composition for Query 6 file

Red: PULocationID

Light blue: list of DOLocationID

Green: sorted list of trip_distance in ascending order relative to DOLocationID